

# Introduction to Data Science - 1MS041

Benny Avelin

Department of Mathematics

HT 2023

# Recall from last time

---

- We defined a stochastic process as a index family of random variables  $X_\alpha$  where  $\alpha \in I$  is the index set. The base example is a sequence of i.i.d. random variables,  $X_1, \dots, X_n$ , here the index set is  $\mathbb{N}$ .

# Recall from last time

---

- We defined a stochastic process as a index family of random variables  $X_\alpha$  where  $\alpha \in I$  is the index set. The base example is a sequence of i.i.d. random variables,  $X_1, \dots, X_n$ , here the index set is  $\mathbb{N}$ .
- We defined the concept of Markov chain, which is a stochastic process which takes a finite number of states and its dependency on the past is only the previous value, i.e.

$$\mathbb{P}(X_{t+1} = x \mid X_1, \dots, X_t) = \mathbb{P}(X_{t+1} = x \mid X_t).$$

## Recall from last time

---

- We defined a homogeneous Markov chain, as one where the transition probabilities does not depend on the index  $t$  (or time).

$$\mathbb{P}(X_{t+1} = y \mid X_t = x) = P_{xy}$$

that is the probability of transitioning from state  $x$  to state  $y$  does not depend on the particular time  $t$ .

# Recall from last time

---

- We defined a homogeneous Markov chain, as one where the transition probabilities does not depend on the index  $t$  (or time).

$$\mathbb{P}(X_{t+1} = y \mid X_t = x) = P_{xy}$$

that is the probability of transitioning from state  $x$  to state  $y$  does not depend on the particular time  $t$ .

- Since  $X_t \in \mathbb{X}$ , where  $\mathbb{X}$  is called the state space and  $\mathbb{X}$  is a finite set. We can define a matrix  $P_{xy}$  where  $x, y \in \mathbb{X}$ . We call this matrix the transition matrix.

# Recall from last time

---

- We defined a homogeneous Markov chain, as one where the transition probabilities does not depend on the index  $t$  (or time).

$$\mathbb{P}(X_{t+1} = y \mid X_t = x) = P_{xy}$$

that is the probability of transitioning from state  $x$  to state  $y$  does not depend on the particular time  $t$ .

- Since  $X_t \in \mathbb{X}$ , where  $\mathbb{X}$  is called the state space and  $\mathbb{X}$  is a finite set. We can define a matrix  $P_{xy}$  where  $x, y \in \mathbb{X}$ . We call this matrix the transition matrix.
- We can use the transition matrix to compute how distributions change when the Markov chain progresses. I.e. let us assume that at time  $t$  we have a distribution over the states  $p_t$ , then the distribution at time  $t + 1$  is computed as

$$p_{t+1} = p_t P \quad p_t = p_0 P^t.$$

# Today

---

Today we will take a look at the pattern recognition problem.


- We will see the first algorithm developed for this problem the "Perceptron", which eventually became neural networks.
- We will study its limitations and how to overcome them by a concept known as kernelization.

# Supervised learning

---

## Setup

1. The generator of the data  $G$
2. The supervisor  $S$
3. The learning machine  $LM$ .





# Pattern recognition

---

- Supervisors conditional distribution  $F(y|x)$  is discrete, and can take  $k$  different values,  $y = 0, \dots, k - 1$ .

# Pattern recognition

---

- Supervisors conditional distribution  $F(y|x)$  is discrete, and can take  $k$  different values,  $y = 0, \dots, k - 1$ .
- Model space  $\mathcal{M} = \{g_\lambda(x) : g_\lambda(x) \in \{0, \dots, k - 1\}\}$ .

# Pattern recognition

---

- Supervisors conditional distribution  $F(y|x)$  is discrete, and can take  $k$  different values,  $y = 0, \dots, k - 1$ .
- Model space  $\mathcal{M} = \{g_\lambda(x) : g_\lambda(x) \in \{0, \dots, k - 1\}\}$ .
- $g_\lambda$  a **decision function, decision rule, classifier**.

# Pattern recognition

---

- Supervisors conditional distribution  $F(y|x)$  is discrete, and can take  $k$  different values,  $y = 0, \dots, k - 1$ .
- Model space  $\mathcal{M} = \{g_\lambda(x) : g_\lambda(x) \in \{0, \dots, k - 1\}\}$ .
- $g_\lambda$  a **decision function, decision rule, classifier**.
- 0 – 1 loss

$$L(z, u) = \begin{cases} 0 & \text{if } y = u \\ 1 & \text{if } y \neq u \end{cases}$$

# Pattern recognition

- Supervisors conditional distribution  $F(y|x)$  is discrete, and can take  $k$  different values,  $y = 0, \dots, k - 1$ .
- Model space  $\mathcal{M} = \{g_\lambda(x) : g_\lambda(x) \in \{0, \dots, k - 1\}\}$ .
- $g_\lambda$  a **decision function, decision rule, classifier**.
- 0 - 1 loss

$$L(z, u) = \begin{cases} 0 & \text{if } y = u \\ 1 & \text{if } y \neq u \end{cases}$$

## The pattern recognition problem

Minimize

$$R(\lambda) = \int L(y, g_\lambda(x)) dF(x, y) = \mathbb{E}[L(Y, g_\lambda(X))]$$

where  $(X, Y) \sim F(x, y)$ , where  $g_\lambda \in \mathcal{M}$ .

# Example

## From regression to decision function

The regression problem. The goal is to estimate the function  $\mathbb{E}[Y | X]$ , where  $Y$  are the labels and  $X$  are the features. We assume  $f_X$  is a fixed density, that we will not care about and instead assume that  $f_{Y|X}$  is part of a parametrized family, i.e.  $f_{Y|X} = p_{\alpha^*, X}$  for some  $\alpha^*$ .

- $p_{\alpha^*, X} = N(\alpha_1 X + \alpha_2, \alpha_3^2)$ , Linear regression
- $p_{\alpha^*, X} = \text{Bernoulli}(G(\alpha_1 X + \alpha_2))$ ,

$$G(x) = \frac{1}{1 + e^{-x}}$$

Logistic regression

## Definition

Let  $r(x) = \mathbb{E}[Y \mid X]$ , then the **Bayes classification rule**  $h^*$  is

$$h^*(x) = \begin{cases} 1 & \text{if } r(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

## Definition

Let  $r(x) = \mathbb{E}[Y \mid X]$ , then the **Bayes classification rule**  $h^*$  is

$$h^*(x) = \begin{cases} 1 & \text{if } r(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

Lets say we use Logistic regression and fit the function  $G(\hat{\alpha}_1 X + \hat{\alpha}_2)$ , now the model is that  $p_{\hat{\alpha}, X} = \text{Bernoulli}(G(\hat{\alpha}_1 X + \hat{\alpha}_2))$  which is a model of the conditional density, i.e.

$$\hat{r}(x) = \mathbb{E}_{p_{\hat{\alpha}, X}}[Y] = G(\hat{\alpha}_1 X + \hat{\alpha}_2)$$



## Definition

Let  $r(x) = \mathbb{E}[Y | X]$ , then the **Bayes classification rule**  $h^*$  is

$$h^*(x) = \begin{cases} 1 & \text{if } r(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

Lets say we use Logistic regression and fit the function  $G(\hat{\alpha}_1 X + \hat{\alpha}_2)$ , now the model is that  $p_{\hat{\alpha}, X} = \text{Bernoulli}(G(\hat{\alpha}_1 X + \hat{\alpha}_2))$  which is a model of the conditional density, i.e.

$$\hat{r}(x) = \mathbb{E}_{p_{\hat{\alpha}, X}}[Y] = G(\hat{\alpha}_1 X + \hat{\alpha}_2)$$

thus we can construct

$$g_{\hat{\alpha}}(x) = \begin{cases} 1 & \text{if } \hat{r}(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

## Definition

Let  $r(x) = \mathbb{E}[Y | X]$ , then the **Bayes classification rule**  $h^*$  is

$$h^*(x) = \begin{cases} 1 & \text{if } r(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

Lets say we use Logistic regression and fit the function  $G(\hat{\alpha}_1 X + \hat{\alpha}_2)$ , now the model is that  $p_{\hat{\alpha}, X} = \text{Bernoulli}(G(\hat{\alpha}_1 X + \hat{\alpha}_2))$  which is a model of the conditional density, i.e.

$$\hat{r}(x) = \mathbb{E}_{p_{\hat{\alpha}, X}}[Y] = G(\hat{\alpha}_1 X + \hat{\alpha}_2)$$

thus we can construct

$$g_{\hat{\alpha}}(x) = \begin{cases} 1 & \text{if } \hat{r}(x) > 1/2 \\ 0 & \text{otherwise.} \end{cases}$$

**Note!!**

This is what scikit-learns LogisticRegression predict function does!

# Pattern recognition

---

## Pattern recognition problem from CS perspective

Suppose we have  $n$  training data points  $T_n := ((X_i, Y_i))_{i=1}^n$  and are interested in a classification rule  $h(X)$  that uses  $T_n$  to *predict*, i.e., assign labels to previously unseen data  $X$ .

# Pattern recognition

---

## Pattern recognition problem from CS perspective

Suppose we have  $n$  training data points  $T_n := ((X_i, Y_i))_{i=1}^n$  and are interested in a classification rule  $h(X)$  that uses  $T_n$  to *predict*, i.e., assign labels to previously unseen data  $X$ .

- We want our classification rule  $h$ , which is typically an algorithm, to perform well on previously unseen data by learning from the training data. This is known as *generalization*.

## Definition

The space  $\mathcal{X}$  where  $X_i$  belongs to is called the *instance space* or *feature space* and the space  $\mathcal{Y}$  where  $Y_i$  belongs to is called the *label space*.

## Definition

The space  $\mathcal{X}$  where  $X_i$  belongs to is called the *instance space* or *feature space* and the space  $\mathcal{Y}$  where  $Y_i$  belongs to is called the *label space*.

Typically,  $\mathcal{X}$  is a subset of  $\mathbb{R}^d$  and  $\mathcal{Y}$  is binary label space either as  $\{0, 1\}$  or  $\{-1, 1\}$ . For example,  $\mathcal{X}$  can be  $\{0, 1\}^d$  to indicate the presence or absence of something in the instance space, say a specific set of words in an email if the task is to classify emails with labels 0 and 1 for non-spam or spam.

## Remark

To connect back to our previous terminology, we see that the data space  $\mathbb{X} = (\mathcal{X}, \mathcal{Y})$  (we will later write  $\mathbb{X} \times \mathbb{Y}$  to avoid  $X$  being used for both feature and label) is split into the feature space and the label space. The random variable we are observing is a pair  $(X, Y)$  and a collection of  $n$  samples is the training dataset.

Let us say that we are trying to device a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .




Let us say that we are trying to device a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .

- Simple rules consists of finding a hyperplane that separate out the data, this is called a **linear separator**.

Let us say that we are trying to device a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .

- Simple rules consists of finding a hyperplane that separate out the data, this is called a **linear separator**.



# How do we find such a separator?

---

- In two or three dimensions it is easy by just looking at the image

# How do we find such a separator?

---

- In two or three dimensions it is easy by just looking at the image
- In higher dimension, i.e. when the number of features is high, we have to find an automated way of doing this!

# The Perceptron algorithm

---

- The perceptron is a linear classification rule

# The Perceptron algorithm

---

- The perceptron is a linear classification rule
- The first training algorithm was devised by Frank Rosenblatt at Cornell Aeronautical Laboratory in 1957.

# The Perceptron algorithm

---

- The perceptron is a linear classification rule
- The first training algorithm was devised by Frank Rosenblatt at Cornell Aeronautical Laboratory in 1957.
- The goal was to build a machine that could recognize hand written letters.

# Lets explain the algorithm

---

## The dimension trick

We start by increasing the dimension, i.e. lets say we have the features  $X = (x_1, \dots, x_k)$ , i.e. we have  $k$  dimensions (features), then we can create  $\tilde{X} = (x_1, \dots, x_k, 1)$  which is now in  $k + 1$  dimensions. The point is that we can write

$$X \cdot w + c = 0 \iff \tilde{X} \cdot \tilde{w} = 0$$

where  $\tilde{w} = (w_1, \dots, w_k, c)$ .



# Lets explain the algorithm

---

## The dimension trick

We start by increasing the dimension, i.e. lets say we have the features  $X = (x_1, \dots, x_k)$ , i.e. we have  $k$  dimensions (features), then we can create  $\tilde{X} = (x_1, \dots, x_k, 1)$  which is now in  $k + 1$  dimensions. The point is that we can write

$$X \cdot w + c = 0 \iff \tilde{X} \cdot \tilde{w} = 0$$

where  $\tilde{w} = (w_1, \dots, w_k, c)$ .

## Warning

We will from now on assume that we have done **the dimension trick** but we go back to the notation  $X$  and  $w$ .

# The perceptron algorithm

---

Let us say that we are trying to device a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .

# The perceptron algorithm

---

Let us say that we are trying to devise a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .

In the following we can take a training dataset  $\{(x_1, y_1), \dots, (x_n, y_n)\}$ . For the perceptron the goal is to find

$$(w \cdot x_i)y_i > 0, \quad i = 1, \dots, n.$$

# The perceptron algorithm

---

Let us say that we are trying to devise a classification rule based on instance space  $\mathcal{X} = \mathbb{R}^d$  and label space  $\mathcal{Y} = \{-1, 1\}$ .

In the following we can take a training dataset  $\{(x_1, y_1), \dots, (x_n, y_n)\}$ . For the perceptron the goal is to find

$$(w \cdot x_i)y_i > 0, \quad i = 1, \dots, n.$$

## The perceptron algorithm

1.  $w = 0$
2. while there exists  $x_i$  with  $x_i y_i \cdot w \leq 0$ , update  $w := w + x_i y_i$

# Convergence is guaranteed if linearly separable

---

## Theorem

*If there exists  $w^*$  such that  $w^* \cdot x_i y_i \geq 1$  for all  $i$ . Then the perceptron algorithm finds a  $w$  satisfying  $w \cdot x_i y_i \geq 0$  for all  $i$  in at most  $r^2 |w^*|^2$  updates, where  $r = \max_i |x_i|$ .*

# Convergence is guaranteed if linearly separable

---

## Theorem

*If there exists  $w^*$  such that  $w^* \cdot x_i y_i \geq 1$  for all  $i$ . Then the perceptron algorithm finds a  $w$  satisfying  $w \cdot x_i y_i \geq 0$  for all  $i$  in at most  $r^2 |w^*|^2$  updates, where  $r = \max_i |x_i|$ .*

## Warning

If the set is not linearly separable then the algorithm will not converge, in fact it will go all over the place.

# Kernelization

---

## The idea

If we cannot separate our data, then perhaps we can find a function  $\phi$  which maps  $\mathbb{R}^{k+1} \rightarrow \mathbb{R}^m$  where  $m > k + 1$ , such that our data becomes linearly separable.

# Kernelization

---

## The idea

If we cannot separate our data, then perhaps we can find a function  $\phi$  which maps  $\mathbb{R}^{k+1} \rightarrow \mathbb{R}^m$  where  $m > k + 1$ , such that our data becomes linearly separable.

So if we transform the  $x \rightarrow \phi(x)$  for some good transformation  $\phi$  then our perceptron will try to solve

$$w \cdot \phi(x_i) l_i > 0$$



## The perceptron algorithm

1.  $w = 0$
2. while there exists  $x_i$  with  $x_i y_i \cdot w \leq 0$ , update  $w := w + x_i y_i$

The weight after  $n$  steps will have the form

$$w = \sum_{i=1}^n c_i \phi(x_i)$$

for numbers  $c_i$ . The perceptron algorithm becomes just addition and subtraction of certain  $c_i$ 's by 1.

Furthermore

$$w \cdot \phi(x_i) = \sum_{j=1}^n c_j \phi(x_j) \cdot \phi(x_i) = \sum_{j=1}^n c_j k_{ij}$$

where  $k_{ij} = \phi(x_i) \cdot \phi(x_j)$ .

## The kernelized perceptron algorithm

1.  $c = 0$
2. While there exists an  $i$  such that  $(Kc)_i y_i \leq 0$  update

$$c_i := c_i + y_i.$$

# Kernel trick

---

## Issue

If  $\phi(x_i)$  is high dimensional there is a big computational cost to computing it.

What if we had a function  $k(x, y)$  that could be written as

$$k(x, y) = \phi(x) \cdot \phi(y)$$

for some  $\phi$  and  $k$  is easier to compute, then our life would be simpler. Also, what if we are given a function  $k(x, y)$  and we would like to know if it is a “kernel function”.

# Kernel functions

---

## Definition

We call a function  $k(x, y) : \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}$  a kernel function if there is a mapping  $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^m$  (for some  $m$ ) such that  $k(x, y) = \phi(x) \cdot \phi(y)$ .

- $k(x, y) = (\gamma x \cdot y + r)^k$ , ( $k \in \mathbb{N}$ ) polynomial
- $k(x, y) = x \cdot y$ , linear
- $k(x, y) = e^{-\gamma |x-y|}$ , called Radial Basis Function
- $k(x, y) = \tanh(\gamma x \cdot y + r)$ , sigmoidal